
PREDICTING CUSTOMER CHURN USING MACHINE
LEARNING TO UNCOVER HIDDEN PATTERN

Student Name: Priya dharshini.R

Register Number: 510623243041

**Institution: C.ABDUL HAKEEM COLLEGE OF
ENGINEERING AND TECHNOLOGY**

Department: BTECH AI&DS

Date of Submission: 08/05/2025

Github Repository Link: https://github.com/Priyadharshini-R123/phase_2.git

1. Problem Statement

Customer churn—the loss of clients or subscribers—is a critical issue for subscription-based businesses, as acquiring new customers is often more costly than retaining existing ones. Despite traditional efforts to mitigate churn, many underlying behavioral patterns and risk indicators remain undetected due to the complexity and volume of customer data.

This project aims to develop a machine learning model that can accurately predict customer churn by analyzing historical customer data and uncovering hidden patterns related to engagement, transaction history, demographics, and support interactions. By identifying high-risk customers early, the organization can proactively implement targeted retention strategies, reduce churn rates, and enhance overall customer lifetime value.

2. Project Objectives

To develop a machine learning model that predicts customer churn by analyzing historical customer data and uncovering hidden behavioral patterns, enabling proactive retention strategies and improved customer lifetime value.

3. Flowchart of the Project Workflow

Data Collection

→ *Gather customer data (transactions, demographics, usage patterns, support logs, etc.)*

Data Preprocessing

- *Handle missing values*
- *Encode categorical variables*
- *Normalize/scale numerical data*
- *Feature engineering*

Exploratory Data Analysis (EDA)

- *Visualize churn trends*
- *Identify correlations and hidden patterns*

Feature Selection

- *Select relevant variables using statistical tests or model-based methods*

Model Selection & Training

- *Train ML models (e.g., Logistic Regression, Random Forest, XGBoost, Neural Networks)*
- *Perform cross-validation*

Model Evaluation

- *Evaluate with metrics like accuracy, precision, recall, F1-score, AUC-ROC*

Model Interpretation

- *Use SHAP, LIME, or feature importance to understand key churn drivers*

Deployment

→ *Integrate model into a production system or dashboard for real-time churn prediction*

Monitoring & Updating

→ *Monitor model performance*

→ *Retrain periodically with new data*

4. Data Description

<i>Feature Name</i>	<i>Description</i>
<i>CustomerID</i>	<i>Unique identifier for each customer</i>
<i>Gender</i>	<i>Customer's gender (e.g., Male, Female)</i>
<i>Age</i>	<i>Age of the customer</i>
<i>Tenure</i>	<i>Number of months the customer has been with the company</i>
<i>SubscriptionType</i>	<i>Type of subscription plan (e.g., Basic, Premium, Family)</i>
<i>MonthlyCharges</i>	<i>Amount charged monthly</i>
<i>TotalCharges</i>	<i>Total amount charged to the customer</i>
<i>PaymentMethod</i>	<i>Method used by the customer to pay (e.g., Credit Card, Bank Transfer)</i>

<i>Feature Name</i>	<i>Description</i>
<i>ContractType</i>	<i>Contract duration (e.g., Month-to-month, One year, Two year)</i>
<i>InternetService</i>	<i>Type of internet service (e.g., DSL, Fiber Optic, None)</i>
<i>OnlineSecurity</i>	<i>Whether the customer has online security add-on (Yes/No)</i>
<i>TechSupport</i>	<i>Whether the customer has technical support add-on (Yes/No)</i>
<i>StreamingServices</i>	<i>Use of streaming services (e.g., TV, Movies, Music)</i>
<i>CustomerSupportCalls</i>	<i>Number of times the customer contacted support</i>
<i>LastInteractionDate</i>	<i>Date of the last customer activity</i>
<i>Churn (Target)</i>	<i>Whether the customer churned (1 = Yes, 0 = No)</i>

5. Data Preprocessing

1. Data Cleaning

- **Remove duplicates:** Drop any duplicate rows using `drop_duplicates()`.
- **Handle missing values:**
 - Numerical: Fill with median/mean or use interpolation.

- *Categorical: Fill with mode or use a placeholder like "Unknown".*

2. Feature Engineering

- *Create new features (e.g., $AverageMonthlySpend = TotalCharges / Tenure$).*
- *Convert dates: Calculate recency ($days_since_last_interaction$ from $LastInteractionDate$).*
- *Bin continuous variables: Group age or tenure into categories if helpful.*

3. Encoding Categorical Variables

- *Label Encoding for binary features (e.g., Yes/No).*
- *One-Hot Encoding for multi-class features (e.g., $SubscriptionType$, $PaymentMethod$).*

4. Scaling / Normalization

- *Use **StandardScaler** or **MinMaxScaler** for continuous features (e.g., $MonthlyCharges$, $TotalCharges$) to normalize the range.*

5. Outlier Detection and Treatment

- *Use Z-score or IQR methods to detect and optionally cap/transform outliers in numerical columns.*

6. Class Imbalance Handling (if churned vs. not churned is imbalanced)

- *Use **SMOTE** (Synthetic Minority Over-sampling Technique) to oversample minority class.*
- *Or use **class weights** in model training.*

7. *Train-Test Split*

- *Split the dataset (e.g., 80% training, 20% test) using stratified sampling to maintain churn ratios.*

8. *Save Preprocessed Data*

- *Optionally, export the cleaned and encoded dataset to CSV or a binary format (e.g., pickle) for reuse.*

6. Exploratory Data Analysis (EDA)

1. *Understand the Target Variable*

- *Value Counts & Distribution*

python

CopyEdit

```
df['Churn'].value_counts(normalize=True).plot(kind='bar', title='Churn Distribution')
```

- *Helps identify class imbalance.*

2. *Summary Statistics*

- *Use `df.describe()` to summarize numerical features.*
- *Use `df.info()` to check data types and null values.*

3. *Univariate Analysis*

- *Numerical Features: Histograms or KDE plots (e.g., Age, MonthlyCharges, Tenure)*

python

CopyEdit

```
df['Tenure'].hist(bins=20)
```

- **Categorical Features:** Bar plots

python

CopyEdit

```
df['ContractType'].value_counts().plot(kind='bar')
```

4. Bivariate Analysis

- **Churn vs. Numerical Variables:** Boxplots or violin plots

python

CopyEdit

```
sns.boxplot(x='Churn', y='MonthlyCharges', data=df)
```

- **Churn vs. Categorical Variables:** Stacked bar plots or heatmaps

python

CopyEdit

```
pd.crosstab(df['ContractType'], df['Churn']).plot(kind='bar', stacked=True)
```

5. Correlation Analysis

- **Correlation Heatmap** for numerical variables

python

CopyEdit

```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
```

- Reveals multicollinearity or redundant features.

6. Feature Interactions

- Use pair plots or scatter matrix to explore interactions (use sample to reduce overload)

python

CopyEdit

```
sns.pairplot(df[['Tenure', 'MonthlyCharges', 'TotalCharges', 'Churn']],  
hue='Churn')
```

7. Customer Segmentation (Optional)

- Apply clustering (e.g., KMeans) or PCA for dimensionality reduction to visualize customer groups.

7. Feature Engineering

1. Basic Derived Features

- $AverageMonthlySpend = TotalCharges / Tenure$
(Use only if $Tenure > 0$, avoid divide-by-zero)
- $IsSenior = Age > 60$
- $HasMultipleServices = Combine\ streaming,\ internet,\ and\ security\ flags$

2. Customer Engagement Features

- $EngagementScore = Weighted\ score\ of\ usage/activity-based\ features\ (e.g.,\ usage\ of\ streaming,\ support\ calls,\ internet\ type)$
- $DaysSinceLastInteraction = Current\ date - LastInteractionDate$

- $SupportCallRate = CustomerSupportCalls / Tenure$

3. Subscription Characteristics

- $ContractLength = Map \text{ "Month-to-month" } \rightarrow 1, \text{ "One year" } \rightarrow 12, \text{ "Two year" } \rightarrow 24$
- $IsAutoPay = PaymentMethod \text{ in } ['Credit card (automatic)', 'Bank transfer (automatic)']$
- $IsLongTermCustomer = Tenure > 24 \text{ (arbitrary threshold)}$

4. Binary Flags for Add-ons

- $UsesTechSupport = Yes/No \text{ to } 1/0$
- $UsesStreaming = Combine \text{ multiple streaming services (TV, music, etc.)}$
- $HasOnlineSecurity = Binary \text{ flag for security service}$

5. Customer Lifetime Metrics

- $LifetimeValue = MonthlyCharges \times Tenure$
- $ChurnRiskScore = Create \text{ rule-based score using a weighted sum of red flags (e.g., short tenure, high support calls, month-to-month contract)}$

6. Interaction Features

- $MonthlyCharges \times ContractLength \rightarrow Measures \text{ financial commitment}$
- $Tenure \times ContractLength \rightarrow Loyalty-adjusted \text{ commitment}$

Tips for Good Feature Engineering

- Consider business logic: Features must make real-world sense.
- Avoid data leakage: Don't use features that are only available after churn.
- Evaluate feature importance later using tree-based models or SHAP values.

8. Model Building

1. Split Data

python

CopyEdit

from sklearn.model_selection import train_test_split

```
X = df.drop('Churn', axis=1)
```

```
y = df['Churn']
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, stratify=y, random_state=42  
)
```

2. Select and Train Models

Try several models to compare performance.

python

CopyEdit

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from xgboost import XGBClassifier
```

```
models = {
```

```
    "Logistic Regression": LogisticRegression(max_iter=1000),
```

```
    "Random Forest": RandomForestClassifier(n_estimators=100),
```

```
"XGBoost": XGBClassifier(use_label_encoder=False, eval_metric='logloss')  
}
```

```
for name, model in models.items():
```

```
    model.fit(X_train, y_train)
```

```
    print(f'{name} trained.')  
  


---


```

3. Evaluate Models

Use appropriate metrics for binary classification.

python

CopyEdit

```
from sklearn.metrics import accuracy_score, precision_score, recall_score,  
f1_score, roc_auc_score
```

```
def evaluate_model(model, X_test, y_test):
```

```
    y_pred = model.predict(X_test)
```

```
    y_prob = model.predict_proba(X_test)[:, 1]
```

```
    return {
```

```
        "Accuracy": accuracy_score(y_test, y_pred),
```

```
        "Precision": precision_score(y_test, y_pred),
```

```
"Recall": recall_score(y_test, y_pred),  
"F1": f1_score(y_test, y_pred),  
"ROC AUC": roc_auc_score(y_test, y_prob)  
}
```

```
for name, model in models.items():  
    scores = evaluate_model(model, X_test, y_test)  
    print(f"\n{name} Scores:")  
    for metric, score in scores.items():  
        print(f"{metric}: {score:.4f}")
```

4. Tune the Best Model (Optional)

Use grid search or randomized search for hyperparameter optimization.

python

CopyEdit

```
from sklearn.model_selection import GridSearchCV
```

```
param_grid = {  
    'n_estimators': [100, 200],
```

```
'max_depth': [4, 6, 8],  
}
```

```
grid = GridSearchCV(RandomForestClassifier(), param_grid, cv=5,  
scoring='roc_auc')  
grid.fit(X_train, y_train)
```

```
best_model = grid.best_estimator_
```

5. Interpret the Model

Use feature importance or SHAP to understand churn drivers.

python

CopyEdit

```
import shap
```

```
explainer = shap.TreeExplainer(best_model)
```

```
shap_values = explainer.shap_values(X_test)
```

```
shap.summary_plot(shap_values, X_test)
```

9. Visualization of Results & Model Insights

1. Churn Distribution

Purpose: Understand class imbalance.

python

CopyEdit

```
df['Churn'].value_counts().plot(kind='bar', title='Churn Distribution')
```

✓ 2. Feature Importance (Tree-Based Models)

Purpose: Identify top predictors of churn.

python

CopyEdit

```
importances = best_model.feature_importances_
```

```
features = X_train.columns
```

Plot

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
feat_imp = pd.Series(importances, index=features).sort_values(ascending=False)
```

```
sns.barplot(x=feat_imp[:10], y=feat_imp.index[:10])  
  
plt.title("Top 10 Feature Importances")
```

✓ 3. SHAP Summary Plot

Purpose: Explain model predictions globally and locally.

python

CopyEdit

```
import shap
```

```
explainer = shap.TreeExplainer(best_model)
```

```
shap_values = explainer.shap_values(X_test)
```

```
shap.summary_plot(shap_values, X_test)
```

- **Color:** Red = high feature value, Blue = low feature value
 - **X-axis:** Impact on churn probability
-

✓ 4. Confusion Matrix

Purpose: Visualize model performance.

python

CopyEdit

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

```
ConfusionMatrixDisplay.from_estimator(best_model, X_test, y_test)
```

✓ 5. ROC Curve

Purpose: Evaluate classification threshold and trade-offs.

python

CopyEdit

```
from sklearn.metrics import roc_curve, auc
```

```
y_prob = best_model.predict_proba(X_test)[:, 1]
```

```
fpr, tpr, _ = roc_curve(y_test, y_prob)
```

```
plt.plot(fpr, tpr, label=f"AUC = {auc(fpr, tpr):.2f}")
```

```
plt.plot([0, 1], [0, 1], '--')
```

```
plt.xlabel("False Positive Rate")
```

```
plt.ylabel("True Positive Rate")
```

```
plt.title("ROC Curve")
```

```
plt.legend()
```

✓ 6. Partial Dependence Plot (Optional)

Purpose: Show marginal effect of features.

python

CopyEdit

from sklearn.inspection import plot_partial_dependence

plot_partial_dependence(best_model, X_test, ['Tenure', 'MonthlyCharges'])

🔗 Model Insights (Example)

- **Top Churn Drivers:**
 - Short tenure
 - Month-to-month contracts
 - High monthly charges
 - Multiple support calls
- **Customer Segments at Risk:**
 - New users (<6 months) on flexible contracts
 - Users not using add-ons like online security or tech support
- **Business Actions:**

- *Offer loyalty discounts*
- *Encourage annual plans*
- *Improve support experience*

10. Tools and Technologies Used

Data Collection & Storage

- *CSV / Excel / SQL — Storing and retrieving customer data*
- *Pandas — Data manipulation and analysis*
- *NumPy — Numerical operations*

Data Preprocessing

- *Scikit-learn (sklearn) — Imputation, encoding, scaling, train-test split*
- *Pandas Profiling / Sweetviz — Automated EDA reports (optional)*

Exploratory Data Analysis (EDA)

- *Matplotlib / Seaborn — Visualization (bar plots, histograms, heatmaps)*
- *Plotly — Interactive charts (optional)*

Model Building

- *Scikit-learn — Logistic Regression, Random Forest, model evaluation*
- *XGBoost / LightGBM — Gradient boosting models for high accuracy*

- **Imbalanced-learn** — SMOTE, under-/oversampling

Model Interpretation

- **SHAP** — Explainable AI to interpret feature contributions
- **LIME** — Local model interpretability (optional)
- **Feature Importance Plots** — Built into tree-based models

Model Evaluation

- **Scikit-learn** — ROC curve, confusion matrix, F1 score, accuracy, precision, recall

Model Deployment (Optional)

- **Flask / FastAPI** — Build REST API for model inference
- **Streamlit / Dash** — Create interactive dashboards
- **Docker** — Containerize the model
- **AWS / Azure / GCP** — Deploy model to the cloud

Version Control & Collaboration

- **Git / GitHub** — Versioning and team collaboration
- **Jupyter Notebook** — For prototyping and documentation

11. Team Members and Contributions

NAME	ROLE	RESPONSIBILITY
PRIYADHARSHINI R	Lead	Oversee project development, coordinate team activities, ensure timely delivery of milestones, and contribute to documentation and Data Engineer final
NANDHITHA M	Data Engineer	Collect data from APIs (e.g., Twitter), manage dataset storage, clean and preprocess text data, and ensure quality of input data
Varshini.S, Vaishnavi.A	NLP Specialist / Data	Build sentiment and emotion classification models, perform feature engineering, and evaluate model performance using suitable metrics
Sonika.R	Data Analyst / Visualization	Conduct exploratory data analysis (EDA), generate insights, and develop visualizations such as word clouds, emotion trends, and sentiment



--	--	--